



FREE PLAYBOOK — FOR AGENCY OWNERS & IMOS

The \$3M/Month Ad Playbook

7 strategies I use to scale insurance lead generation past
\$100K/month in ad spend — profitably.

\$3M+/mo managed

1M+ leads generated

\$50M+ premium placed

3,700+ clients served

Jeromy Kovatana — Founder, LeadsBakery
jeromykovatana.com



READ THIS FIRST

Why agencies stall — and what this playbook actually is

Most insurance agencies don't have an advertising problem. They have a **system** problem.

They can get to \$20—30K/month in spend on hustle and a couple of decent ads. Then everything wobbles: CPL creeps up, creative burns out, agents complain about lead quality, and the owner starts making panic decisions inside Ads Manager at 11pm. Sound familiar?

I came into this industry through a side door. Oklahoma State, then Big 4 accounting — Ernst & Young, Grant Thornton — then a CFO seat. Then a bankruptcy that stripped everything down to faith and family and forced a rebuild. When I found insurance lead generation, I realized my finance background was an unfair advantage: **I read ad accounts the way a CFO reads a P&L**. Cost per acquisition isn't a marketing metric to me. It's a line item that decides whether your agency survives.

Today my companies manage over \$3 million per month in ad spend across final expense, veterans programs, mortgage protection, IUL, trucker insurance, and Spanish-language funnels. Everything in this playbook is pulled from those live accounts — not theory, not recycled guru content. These are the seven structural decisions that separate accounts that scale from accounts that burn.

HOW TO USE IT

- **Each strategy is one page.** Why it works, how to run it, and the mistake that kills it.
- **They compound in order.** Economics first, then creative systems, then structure, then follow-up.
- **Score yourself honestly.** Most agencies I audit are running two of the seven. The gap is the opportunity.

A note on scale: nothing here requires \$100K/month to start. It requires deciding to operate like the agencies that get there.

Run Your Ad Account Like a P&L

Scaling dies the moment CPA becomes a "marketing number" instead of a financial one.

WHY IT WORKS

Every vertical has an **allowable CPA** — the most you can pay for a lead and still make money after placement rate, average premium, comp level, and overhead. Most agencies have never calculated it. They judge weeks by "leads felt expensive" instead of by contribution margin. The agencies that scale know their number to the dollar and force the auction to respect it.

HOW TO RUN IT

- **Calculate allowable CPA:** average annual premium × placement rate × your comp – cost to work the lead. That's the ceiling.
Set your target 20—30% below it.
- **Use cost cap bidding on every scaling campaign.** Across the \$3M+/month I manage, virtually every dollar runs under a cost cap tied to allowable CPA — not to last week's average.
- **Review weekly like a CFO:** spend, CPL, contact rate, placement, premium per lead source. Volume can lie to you.
Contribution margin can't.
- **Budget by margin, not mood.** Winners earn budget mechanically. Nothing gets scaled because it "feels like it's working."

WATCH OUT

Lowest-cost bidding ("just spend the budget") feels cheaper — until lead quality collapses and your agents stop answering the CRM. Caps exist to make Meta find your buyer, not the cheapest click.

The Weekly Winners Refresh

Creative fatigue is the silent killer above \$3—5K/day. The fix isn't new ads — it's a refresh system.

WHY IT WORKS

At real spend, even great creative wears out in weeks. Frequency creeps, CPL climbs 20—40%, and most buyers respond by either panic-testing junk or white-knuckling a dying campaign. A standing refresh cadence removes the emotion: proven creative gets re-launched into a fresh campaign on a schedule, so the account never depends on one tired ad set.

HOW TO RUN IT

- **Maintain a "Winners" campaign** that holds only creative that has already beaten target CPA at volume.
- **Every 7—10 days, relaunch the winners as a brand-new campaign.** New auction entity, proven assets — you reset fatigue without gambling on unproven creative.
- **Retire ruthlessly.** Anything that can't hold CPA through two consecutive refreshes leaves the roster.
- **Keep a bench.** Your testing pipeline (Strategy 03) should always have 2—3 graduates ready, so a refresh never launches thin.

FROM MY OWN ACCOUNTS

Pull up my largest final expense account and you'll see the cadence in the campaign names themselves: Winners launched **5/11** → **5/18** → **5/26** → **6/04**. Week after week, like clockwork. That rhythm — not any single ad — is what holds CPL flat at scale.

WATCH OUT

Don't refresh losers. A relaunch amplifies whatever you feed it — feed it your best or don't feed it at all.

Testing Grounds, Not Testing Guesses

Testing inside your scaling campaigns poisons the signal in both directions.

WHY IT WORKS

Drop raw creative into a scaling campaign and you tax your winners' delivery while giving the new ad a distorted read. Separate the jobs: scaling campaigns exist to spend efficiently; a dedicated testing campaign exists to buy information. When testing has its own budget and its own scoreboard, every dollar teaches you something — and your scaling campaigns stay clean.

HOW TO RUN IT

- **Run an always-on Testing Ground per vertical** with a fixed daily budget — in my accounts these run \$1.5—2.5K/day, roughly 10—15% of vertical spend.
- **Use dynamic creative** to let Meta assemble hook × body × format combinations you'd never hand-build.
- **Set a graduation rule and honor it:** a test only moves to Winners after beating target CPA at meaningful volume (think 50+ leads, not 5).
- **Feed it weekly:** new hooks, new angles, new formats. The testing pipeline is the supply chain for Strategy 02.

WATCH OUT

Testing budget is a fixed cost of scaling, not a luxury you cut in a bad week. Cut it and four weeks later you're refreshing into an empty bench — at exactly the moment you need fresh winners most.

One Vertical, One Account

Blended accounts produce blended signal — and concentrated risk.

WHY IT WORKS

A final expense buyer, a veteran's spouse, and a 35-year-old IUL prospect are different humans with different economics. Run them through one account and the pixel learns an average customer who doesn't exist — while one compliance flag threatens every vertical you operate at once. Separation keeps the machine learning sharp and the blast radius small.

HOW TO RUN IT

- **One ad account per vertical.** This is exactly how I structure my largest client's operation: separate accounts for final expense, veterans, mortgage protection, IUL, and trucker — each with its own pixel history, its own allowable CPA, its own creative library.
- **One page identity per audience** (more on this in Strategy 06) — the page is part of the targeting.
- **Build compliant from day one.** Insurance runs under Meta's financial products & services special ad category. Structure for it; don't fight it.
- **Contain failures.** When an account gets reviewed or restricted — and at scale, eventually one will — the other verticals keep spending.

WATCH OUT

Don't fragment too early. Prove one vertical end-to-end — stable CPA, real placement data — before you open account number two. Structure follows proof.

Split Creative by Format

A blended campaign will hide your real winner behind your cheapest impression.

WHY IT WORKS

Meta optimizes toward whatever clears the auction cheapest — which is not the same as what produces leads your agents can place. Mix video and static in one campaign and the platform quietly starves one format before you ever get a clean read. Splitting them turns "which creative works" from a guess into a ledger.

HOW TO RUN IT

- **Run video-only and image-only campaigns side by side** in the same account. In my accounts they're literally labeled VID-only and IMG-only — same offer, separated formats, separated budgets.
- **Judge two layers deep:** CPL per format, then contact rate and placement rate per format downstream in the CRM. Cheap leads that never answer aren't cheap.
- **Let budget follow placed premium,** not surface CPL. The format that wins the P&L gets the next dollar.
- **Re-test the split per vertical.** What wins in final expense is not what wins in IUL — different buyer, different age band, different scroll behavior.

WATCH OUT

Don't declare a format winner off a week of data. Formats trade the lead in cycles; the ledger over 3—4 refresh cycles is the truth.

Localize the Front End

Relevance is the cheapest CPL lever left on the table.

WHY IT WORKS

The auction is saturated with generic "life insurance" advertisers. You don't out-bid them — you out-relevance them. A retiree in Boise responds differently to "Idaho Life Coverage for Seniors" than to a national brand. A trucker responds to a brand built for truckers. Specificity raises click-through, which lowers CPM, which drops CPL — before you've touched targeting.

HOW TO RUN IT

- **Build state-level page identities.** I run dedicated pages per state for senior life coverage — same backend, localized front end.
- **Give niches their own brand.** Trucker insurance in my portfolio isn't a campaign inside a general account — it's its own brand, own page, own funnel.
- **Go native in language, not translated.** My Spanish-language final expense funnels (own brand, own creative, own forms) are built in Spanish from scratch — translation kills trust.
- **Match the chain:** ad → page identity → landing page → form should all speak to the same person in the same voice.

WATCH OUT

Every new front end needs its own compliance pass and a seasoning period. Spin up identities faster than you can run them properly and you trade CPL gains for account risk.

Speed-to-Lead, or Don't Bother Scaling

The oldest truth in this industry is still the most violated: a lead worked in five minutes is a different asset than the same lead worked in an hour.

WHY IT WORKS

Scaling ad spend with broken follow-up doesn't grow your agency — it burns cash at a higher rate. Contact rate is the multiplier that sits between CPL and placed premium. Move it from 40% to 65% and you just cut your effective cost per placement by a third without changing a single ad.

HOW TO RUN IT

- **Automate the first 60 seconds:** the instant a lead lands — SMS out, email out, CRM record created, agent assigned. No human in the loop. This stack is why I built TeleCRM and AgenticPro on top of GoHighLevel.
- **Route deliberately:** round-robin or skill-based assignment with accountability — every lead owned by a name within a minute.
- **Run a 14+ day no-answer cadence.** Most placed premium I see comes from leads touched 5+ times. The fortune is in the eighth attempt.
- **Track contact rate like you track CPL.** Put them on the same dashboard; they only mean something together.

WATCH OUT

A good automation stack replaces about three full-time employees of manual work — but only if someone owns it. Unowned automations decay silently while you scale spend on top of them. Assign an owner.

THE PART WHERE MOST PEOPLE STOP

Knowing the system isn't running the system.

Everything in this playbook is simple to understand and brutal to execute consistently — while you're also recruiting agents, managing carriers, and running an agency.

That's the part I do. I take on a small number of agencies and IMOs and build this entire machine — strategy, ad management, creative systems, automation, and the CRM infrastructure underneath it. From wherever you are today to \$100K+/month in profitable spend.

Book a free strategy call → jeromykovatana.com

30 minutes. A real audit of your current setup, a prioritized roadmap, and zero pressure.
Direct booking link: calendar.app.google/UjwwbYsGsF9ZVn6U8

\$3M+/mo managed

1M+ leads generated

\$50M+ premium placed

3× Two Comma Club

© 2026 Jeromy Kovatana. All rights reserved. This playbook is for educational purposes and isn't a guarantee of results — your economics, carriers, and execution will vary.